

Topic- Rate Limiter

What is rate limiter?

→ Control the rate of traffic it could be on basis of anything eg. user_id, IP etc

Why do we need it?

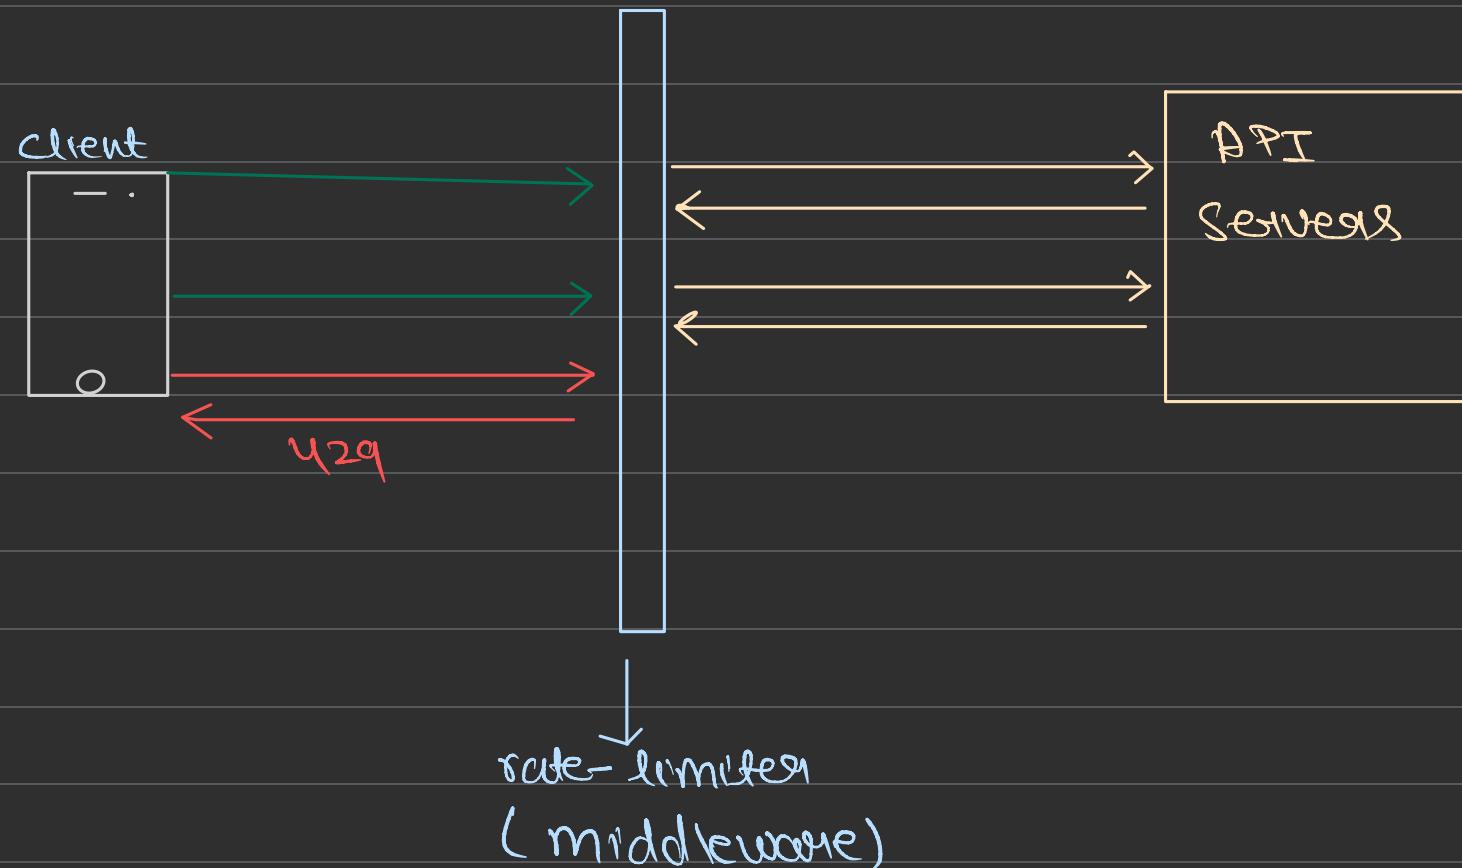
→ Server won't be overloaded.

→ Prevents from DDoS attack.

Example -

every user can't more than 10 requests / sec.

API state limiter.



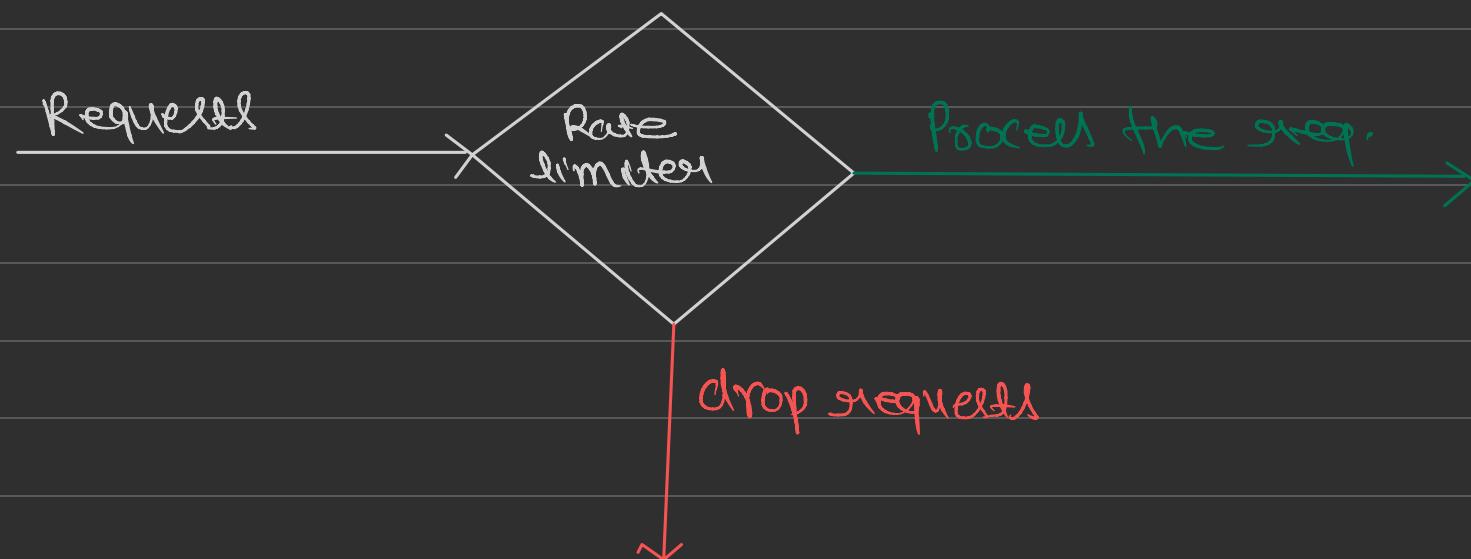
Rate limiter algorithms

Refilling bucket at constant state.

① Token bucket



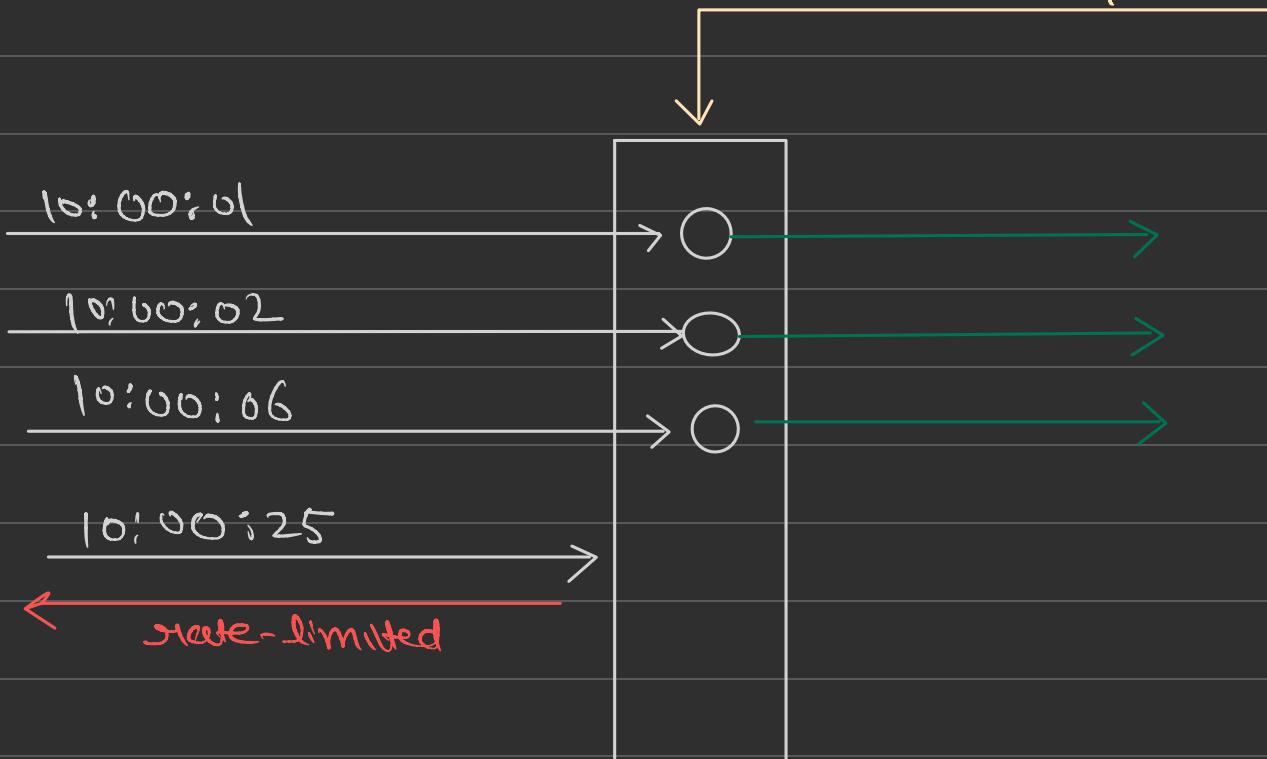
bucket with tokens



e.g. refil rate = 3 tokens / minute

bucket size = 4 tokens.

3 tokens / min.



ISSUE: At the time unit boundaries, we may twice
the traffic.

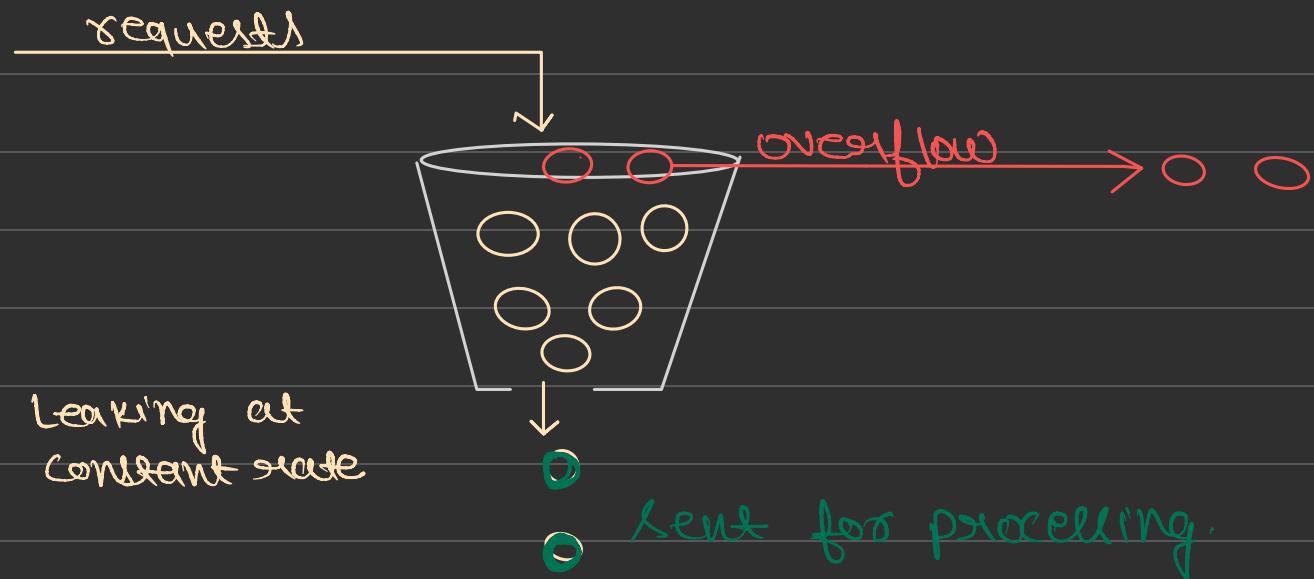
- ① Refill rate decides Avg traffic to your server.
- ② Bucket size signify Peak traffic in a unit of time.

e.g. refill rate = 4 token / min.

Bucket size \geq 16 token.

② Leaking bucket.

Bucket with a given capacity that is leaking at a constant rate. i.e.



Can be implemented via queue, whenever a req.

comes -

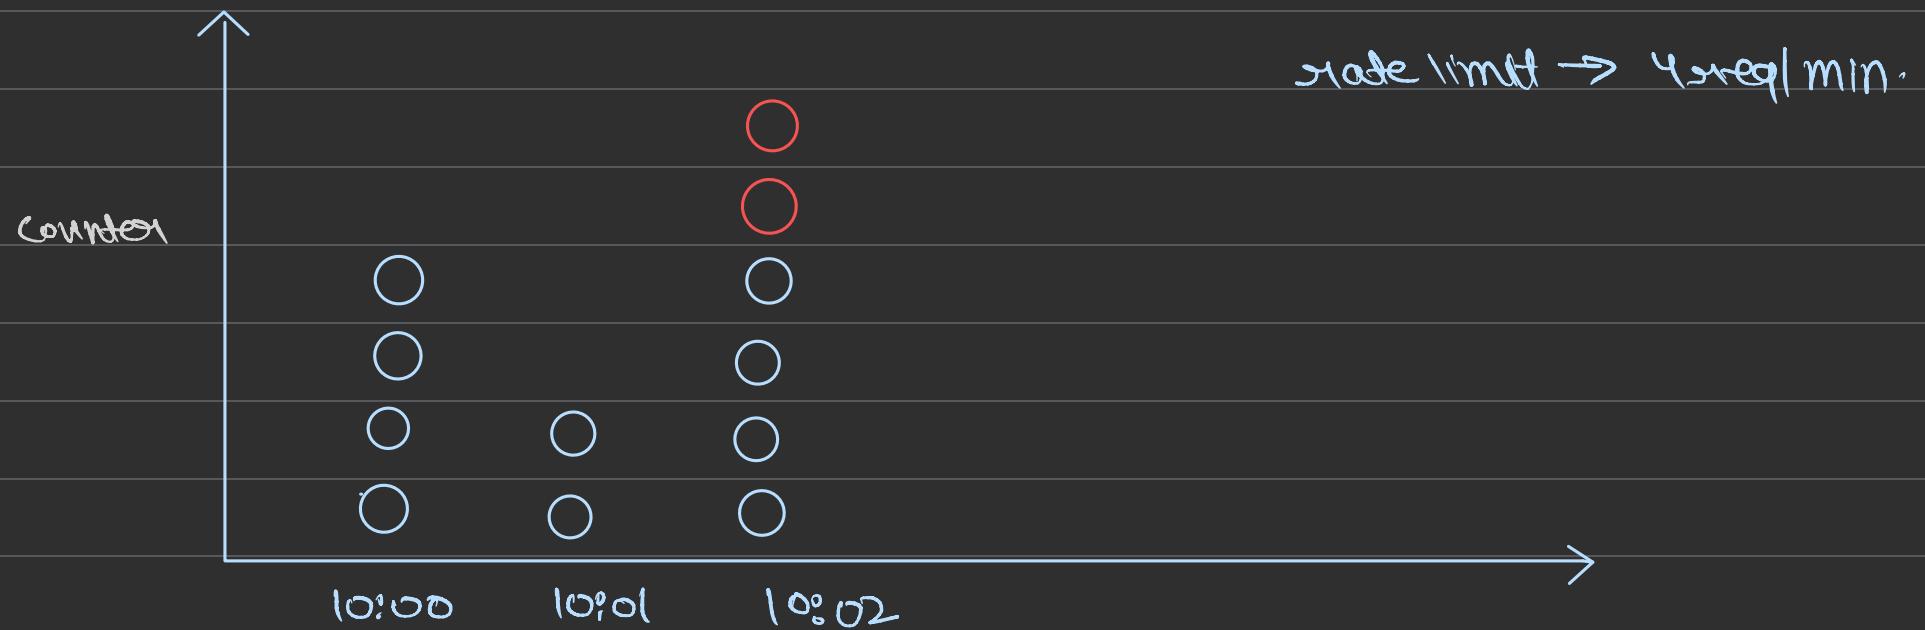
→ If queue isn't full, add them to the queue.
else drop the request.

Also server will keep consuming requests from
the queue at a constant rate.

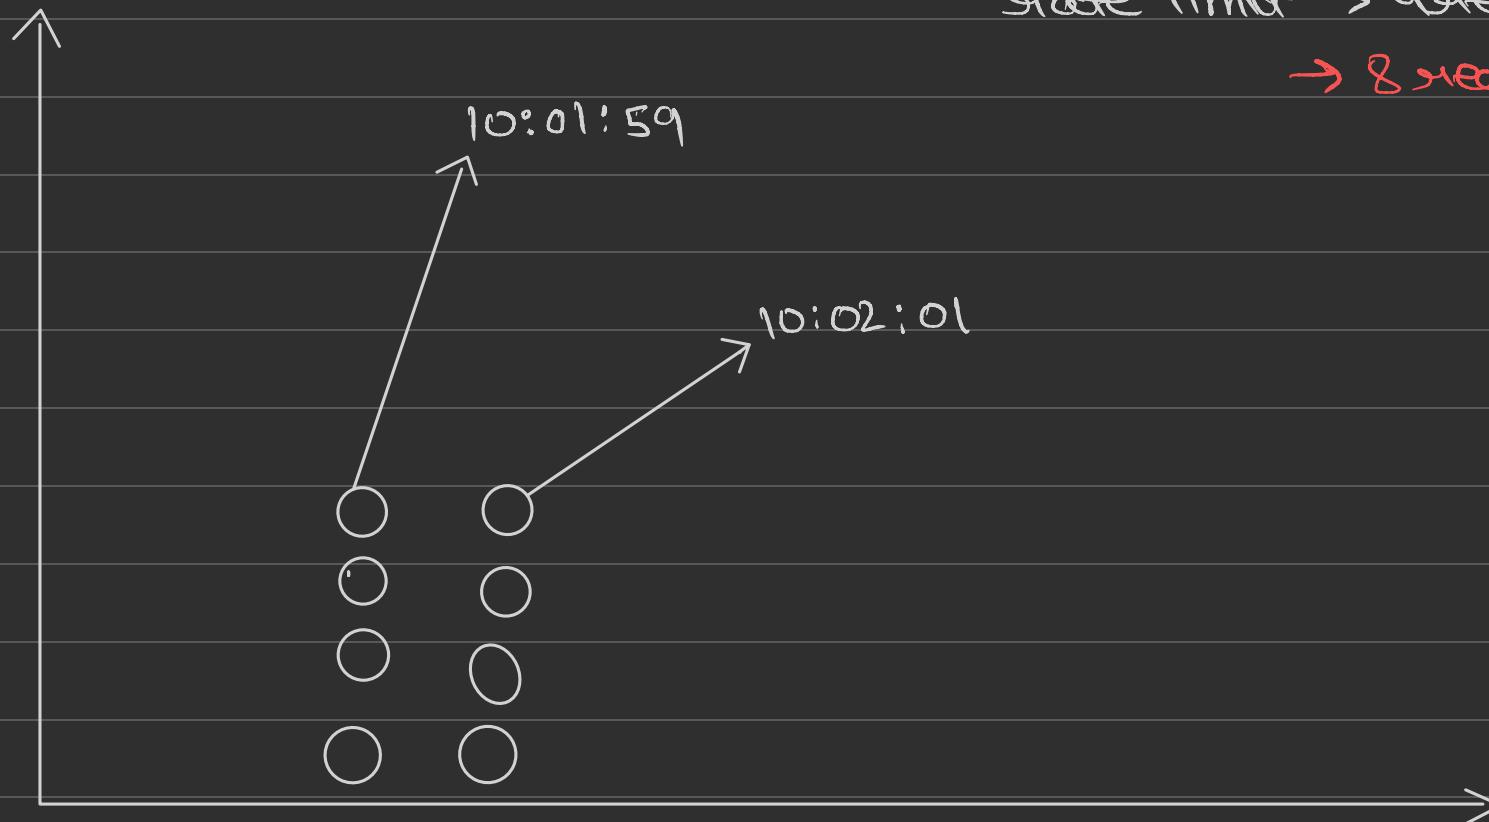
Issue → burst traffic can quickly fill the queue & all
subsequent req. will be rejected.

③ fixed window counter

- ① Create a counter for time window.
- ② When seq come, identify the counter and try to increase the counter, if its less than threshold.



Issue → burst traffic



rate limit → req/min.
→ 8 reqs/min.

⑤ Sliding window log → fixed the issue of static window counter.

- ① keep log of request timestamp.
- ② whenever new req. comes, simply add it in the log.
- ③ remove expired timestamp.
- ④ if size of log is within threshold, then accepts the req. else reject.

Let's take an example.

→ log threshold 3

→ time unit is minute

$\pi_1 \rightarrow 10:00:05$

$\pi_2 \rightarrow 10:00:09$

$\pi_3 \rightarrow 10:00:50$

let say, π_4 comes at $10:01:06$

~~$\pi_1 \rightarrow 10:00:05$~~ expire.

$\pi_2 \rightarrow 10:00:09$

$\pi_3 \rightarrow 10:00:50$

$\pi_4 \rightarrow 10:01:06$

let say, π_5 comes

~~$\pi_1 \rightarrow 10:00:05$~~ expire.

$\pi_2 \rightarrow 10:00:09$

$\pi_3 \rightarrow 10:00:50$

$\pi_4 \rightarrow 10:01:06$

$\pi_5 \rightarrow 10:01:07 \rightarrow \text{drop}$

⑤ Sliding window counter

Combination of sliding window log + fixed window counter.

- ① It always keep track of no. of req. in prev min.
- ② When a req. come, let say at k^{th} fraction of minute. It finds **Approximate traffic** of sliding window.

no. of req. in prev minute \times fraction of prev minute while considering sliding window + no. of req. in current minute.

Let say - state limit = 10 req./min.

no. of req. in prev min (10:01) \rightarrow 9 req.

third req. of current min comes at \rightarrow 10:02:20

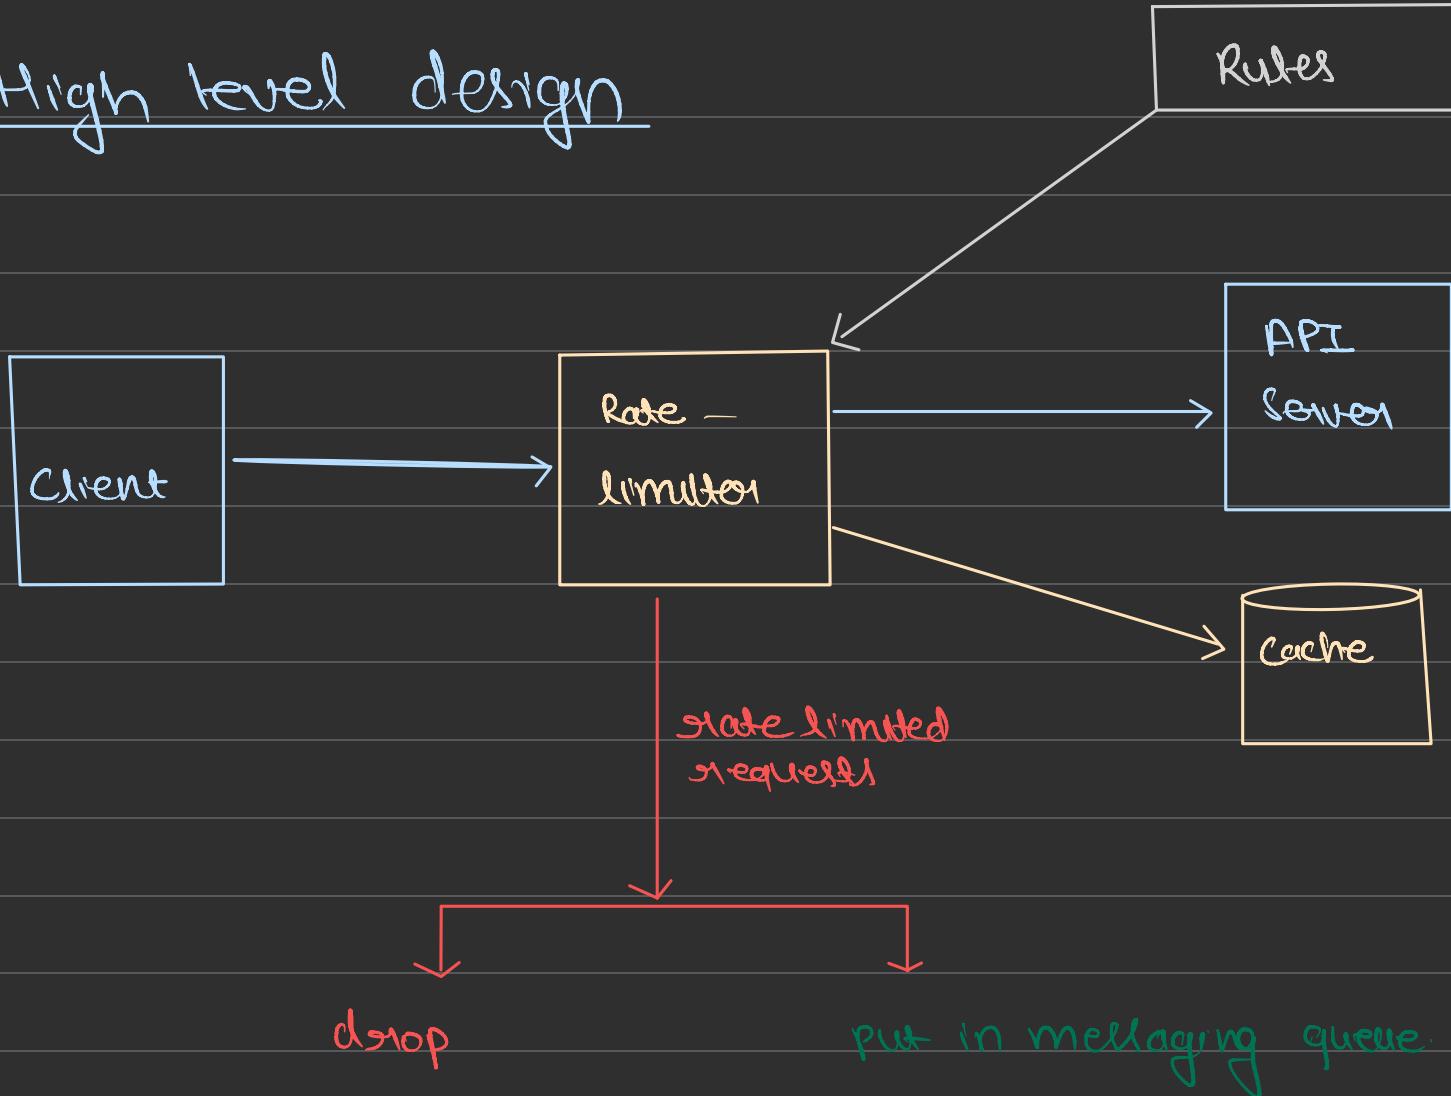
fraction of prev min in current sliding window = $\frac{2}{3}$

so total approximate req. in current sliding window

$$= 9 \times \frac{2}{3} + 2 + 1 = 9$$



High level design



Challenges with distributed rate-limiter.

- ① Race condition \rightarrow Take locks while inc the counter.

let say $x=3$, two concurrent thread execute $x=x+1$

we may end up with $x=4$, instead of 5

② Synchronization

Let say if we have many instance of rate limiter running. challenge is to have a same delta stored on each instance.

→ centralized cache (Redis)

γ1 -

10:00:57

c=1

10/mu.

γ2 -

c=2

γ3 -

c=3

γ4 -

10:01:57 = c=4 c=1

=

-

=

10:02:56

→ ⑨

c=10

10:02:57 → c=1

every rate-limiter host have its own cache

